

rtl-mosaic: A Reuse-Aware Benchmark for LLM-Driven RTL Design

Leonardo Ferreira, Matthew Kotzbauer, Warren Zhu

Team 11, COMPSCI 1440R, Harvard University

{leonardo.ferreira@childrens.harvard.edu, matthewkotzbauer@college.harvard.edu, wzhu@college.harvard.edu}

Abstract—Existing LLM-for-Verilog benchmarks score models on from-scratch generation of RTL code. Real-world chip design tasks, however, tend to be more compositional, with modern SoCs running 60–70% licensed IP by area. We built rtl-mosaic, a system-level harness that decomposes a Verilog spec into subblocks, routes reusable pieces to a curated 30-IP corpus over MCP, and stitches the result through Icarus Verilog testbenches. We then ran 15 frontier LLMs from four providers through the harness on 26 ChipBench problems (9 cpu_ip + 17 self_contain), scoring each against 26 hand-labeled gold entries. Across these runs, IP-reuse skill emerges as a distinct axis from scratch coding: routing F1 spreads from 0.19 to 0.46 across models while scratch pass rate on the same problems stays roughly flat. A reuse-aware metric reveals a model capability that scratch-only benchmarks are unable to capture.

Index Terms—LLM, Verilog, RTL, IP reuse, benchmark, MCP, hardware design

I. INTRODUCTION

A modern SoC is 60–70% licensed IP by area [1]. Companies like Synopsys and Cadence ship thousands of pre-verified IP blocks, and re-authoring every primitive from scratch is incompatible with tape-out economics. Yet existing dominant benchmarks for LLM-driven Verilog (VerilogEval [4], ChipBench [3]) score models exclusively on *scratch generation*, measuring whether a model can produce functionally correct Verilog from a natural-language specification. Neither rewards the agent for recognizing that an off-the-shelf register file would suffice, nor penalizes re-authoring one from gates.

We argue this is the wrong incentive: the capability worth measuring is whether a model recognizes the FIFO is on the shelf and uses it correctly with the right widths and ports. This motivates rtl-mosaic, a system-level harness on top of SiliconMind-V1 [2], and used it as the substrate for a multi-LLM benchmark.

Contributions. We introduce a reuse-aware evaluation protocol for Verilog generation. We ship a 30-IP corpus served over MCP, 26 hand-labeled gold-standard decomposition problems on top of ChipBench, and both a deterministic keyword-overlap router and an embedding-similarity router for comparison. We run 15 frontier LLMs spanning Anthropic, OpenAI, AWS Bedrock (DeepSeek), and Google through the protocol (390 evaluation calls total). The result, summarized in Fig. 1, separates routing skill from scratch skill: F1 spread of 0.27 across models against a much narrower scratch-pass-rate spread on the same problems. We open-source the corpus,




Fig. 1. (a) The rtl-mosaic evaluation pipeline. A spec goes to an LLM-as-planner that emits subblocks tagged REUSE_IP or GENERATE. Reusable subblocks resolve through a deterministic router over a 30-IP corpus served via MCP. Custom logic goes through codegen. The integrator stitches everything into a `TopModule` verified by Icarus Verilog. (b) Each dot is one of 15 frontier LLMs we evaluated. The four colored dots are best-of-provider models. Routing F1 (y-axis, vs. hand-labeled gold) and scratch Verilog pass rate on the same RISC-V CPU IP problems (x-axis) show no diagonal trend, so a model’s scratch coding score gives little information about its IP-routing score.

gold labels, and runner at <https://github.com/MattKotzbauer/rtl-mosaic>.

II. APPROACH

A. The pipeline

For each (LLM, problem) pair, we run the four-stage pipeline shown in Fig. 1(a). The *planner* receives the ChipBench spec and a fixed system prompt instructing it to return JSON: a list of one to six subblocks, each tagged REUSE_IP or GENERATE, with a port specification and a free-text `search_query`. The *router* is fully deterministic and, for every REUSE_IP subblock, queries our 30-IP corpus over an MCP server and selects the IP with maximum keyword overlap against the planner’s search query. Subblocks tagged GENERATE are passed to a separate *codegen* call. The *integrator* concatenates the corpus IP source files with the generated logic and asks an LLM to write the `TopModule`

wiring. The result is compiled by `iverilog -g2012` and verified against the ChipBench testbench.

B. What we score

The benchmark in this paper isolates the *routing* stage. We hold the corpus, the router, and the gold labels fixed across all 15 LLMs. The only thing that changes is the planner, so any difference in the score is a property of the planner LLM. For every problem we have hand-labeled gold consisting of (a) a set of corpus IP IDs that a sensible planner-plus-router should surface, and (b) a per-role decision of whether each subblock should be `REUSE_IP` or `GENERATE`.

We track three metrics: precision (the fraction of router-picked IPs also in the gold IP set), recall (the fraction of gold IPs the router picked), and kind agreement (the per-subblock fraction where the planner’s `REUSE_IP/GENERATE` decision matches the gold label, computed via best-effort name matching between planner subblock names and gold role labels). F1 is our headline number because a model that picks one correct IP and ignores everything else scores a misleadingly high precision.

C. The corpus and gold labels

The corpus is 30 IP modules across six categories: memory (sync/async FIFO, single/dual-port RAM, register file), datapath (32-bit ALU, 16x16 multiplier, 8-bit divider, ripple/CLA adder, subtractor, comparator, barrel shifter), counter (up, down, up/down, BCD, gray-code, Johnson, frequency divider), sequence (LFSR, edge detector, shift register, sign-extend), primitive (T flip-flop, 2/4/8-input mux, 3-to-8 decoder, priority encoder), and clock (frequency divider). Every IP ships with a parameterizable Verilog implementation and an Icarus self-test that we run as part of CI. The set was sized to cover the architectural shapes that appear in ChipBench’s hierarchical `not_self_contain` and `cpu_ip` sets, plus seven additions from the `self_contain` set (multiplexer, dual-port RAM, up/down counter, Johnson counter, gray counter, fractional frequency divider, pipelined multiplier). All IPs ship under MIT.

The gold labels live in `eval/gold_labels.py` as a Python dict, written by hand by one of us after reading each problem’s specification. Each entry records the expected corpus IP IDs, the per-role kind decisions, and a short rationale. The labels reference IP IDs from the full 30-IP corpus, so labels written before a corpus expansion remain stable.

III. IMPLEMENTATION

A. Models tested

We ran 15 LLMs across four providers. From Anthropic via the `claude` CLI: Opus 4.7, Sonnet 4.6, Haiku 4.5. From OpenAI via the `chat-completions` API: GPT-5.4, 5.2, 5.1, 5, and 4.1. From AWS Bedrock in `us-east-1` via the `converse` API: DeepSeek-R1 (through the `us.deepseek.r1-v1:0` inference profile) and DeepSeek V3.2. From Google Gemini via the `generativelanguage.googleapis.com/v1beta`

REST endpoint: 2.5 Pro, 2.5 Flash, 2.5 Flash-Lite, 3 Flash Preview, and 3.1 Flash-Lite Preview. (Three additional Gemini variants (3 Pro Preview, 3.1 Pro Preview) returned 503s on the majority of calls during our testing and were excluded.)

All 15 models receive the same planner system prompt with no few-shot examples, temperature set to zero where the API supports it, and are evaluated against the same 30-IP corpus on the same 26 problems (9 from ChipBench `cpu_ip` and 17 from `self_contain`). We use 16-way thread-pool concurrency with retries handled by the underlying SDKs.

For the headline numbers in Fig. 1(b) and Table I we report the four best-of-provider models with full $n = 26$ coverage: Claude Opus 4.7, GPT-5.2, DeepSeek V3.2, and Gemini 2.5 Flash-Lite. We expanded the eval set from 16 to 26 problems on 2026-05-05 to broaden statistical coverage, and numbers in this report are from the expanded set.

B. Eval scaffolding

The full multi-LLM routing eval is contained in (`eval/multi_routing.py`, ~160 lines). It loads the gold labels, fans out 240 calls across providers, records planner output and router decisions per call, and aggregates into per-(LLM, problem) records and a per-LLM summary. Per-call latency, raw planner output, parsed subblocks, and router decisions are all preserved in JSON for post-hoc analysis. The figure pipeline (`eval/make_figures.py`) regenerates ten plots from the JSON without re-running the eval.

C. Cost and time

A single full sweep across all 15 models runs in under six minutes using 16-way concurrency. The most expensive provider per call was OpenAI (~\$0.03 for GPT-5.4 on a typical planner call, dominated by output tokens). DeepSeek-R1 on Bedrock was the cheapest at roughly \$0.30 for the entire 9-problem run including reasoning-trace tokens. The full reproducible eval, including all per-provider sweeps during development, costed under \$25 total.

IV. RESULTS

A. Routing quality is its own axis

Fig. 1(b) plots each of the 15 LLMs as a dot, with the four colored dots representing best-of-provider headline models. The x-axis is scratch pass rate on the `cpu_ip` set (single-shot Verilog through the ChipBench testbench) and the y-axis is routing F1 against our hand-labeled gold on the 26-problem eval set. The scatter shows no diagonal correlation between the two axes, meaning routing skill does not simply follow overall model strength.

The F1 spread across all 15 models is 0.27 (0.19–0.46) on the 26-problem set, while the pass-rate spread on `cpu_ip` is only 0.11 (most models pass exactly 1/9). On a single-axis benchmark the models would be indistinguishable, whereas on `rtl-mosaic` they separate cleanly. The multi-seed columns of Table I re-rank Opus 4.7 above DeepSeek V3.2 and GPT-5.2, and the per-problem σ shows that single-seed numbers overstate the precision available from one run.

TABLE I

HEADLINE ROUTING F1, BEST-OF-PROVIDER, 26-PROBLEM EVAL (9 CPU_IP + 17 SELF_CONTAIN). SINGLE-SEED COLUMN FROM THE FULL 15-MODEL SWEEP (23–26 SUCCESSFUL CALLS PER MODEL DEPENDING ON API AVAILABILITY), MULTI-SEED COLUMN (N=3 SEEDS) FROM A FOCUSED RE-RUN ON THE FOUR HEADLINE MODELS. PER-PROBLEM σ QUANTIFIES RUN-TO-RUN VARIANCE AT TEMPERATURE ZERO.

Provider	Model	F1 (1 seed)	F1 (3 seeds)	σ_{prob}
Google	Gemini 2.5 Flash-Lite	0.464	0.514	0.094
Anthropic	Claude Opus 4.7	0.265	0.404	0.161
Bedrock	DeepSeek V3.2	0.355	0.390	0.071
OpenAI	GPT-5.2	0.433	0.381	0.126

B. Provider family does not predict the winner

Table I also undercuts the assumption that the strongest provider wins. DeepSeek V3.2 (Bedrock) ties with Anthropic Claude on kind agreement and beats it on F1 by 0.066. Within OpenAI, GPT-5.2 outperforms the newer GPT-5.4 in our run. Our claim is more limited: holding the planner prompt constant across all models, which is the only way to get apples-to-apples comparisons, surfaces the prompt-sensitivity of each model in different ways. A single planner prompt underestimates each model’s full capability, and it remains the right protocol for a first-cut benchmark.

C. Failure modes

We observed four recurring failure modes across the 390 calls. (1) *Hallucinated picks*: the planner correctly tags a 32-bit ALU as REUSE_IP but emits a search query the keyword router resolves to `priority_encoder`, since the router does keyword overlap and a bad query produces a bad pick. (2) *Over-decomposition*: weaker models split a five-line spec into six GENERATE subblocks and never reuse anything. (3) *Under-decomposition*: GPT-5 returns one big GENERATE block on average (1.0 subblocks per problem), so the router never gets called, which explains why GPT-5 sits at the bottom of our leaderboard despite being a competent code model. (4) *Kind-flip*: at temperature zero, some Anthropic models swing between REUSE_IP and GENERATE on the same subblock across reruns. Claude Opus 4.7 went from F1=0.678 in an early 9-problem run to F1=0.265 in our final 26-problem run. Run-to-run variance is real even at deterministic settings, and we run multi-seed (n=3) averaging on the four headline models to put error bars on Table I.

D. The corpus-disambiguation bottleneck and the embedding router

When we expanded the corpus from 20 to 30 IPs, mean F1 across the keyword-router eval dropped, not because corpus coverage worsened but because of corpus disambiguation. Adding `up_down_counter`, `johnson_counter`, and `gray_counter` alongside the existing `up_counter` created near-duplicate keyword profiles, so a planner emitting `search_query="counter"` resolves to whichever IP keyword matches first rather than the architecturally correct one. Replacing the keyword router with cosine similarity

over OpenAI `text-embedding-3-small` representations of each IP’s name, description, and keywords lifts mean F1 by +0.029 across all 15 models, with 9 of 15 improving by ≥ 0.01 (DeepSeek-R1 +0.128, GPT-5.2 +0.100, GPT-5.4 +0.079). This means some degree of the routing bottleneck lives in the router, not just in the planner LLM, within the measured spread.

V. RELATED WORK

ChipGPT [5], VeriGen [6], RTLCode, and HiVeGen target scratch Verilog generation, with HiVeGen retrieving snippets from code corpora rather than IPs. ChipNeMo [7] domain-adapts an open model on EDA scripts. MAGE [8] multi-agent decomposes per-subtask scratch generation. SiliconMind-V1 [2] multi-agent generates one module at a time and is the substrate we extend. None of these prior works score reuse against an IP catalog.

The benchmark side is similar. VerilogEval [4] is single-module, so reuse is degenerate. RealBench targets full SoCs but does not separately score IP retrieval. ChipBench [3] ships an `cpu_ip` subset that is hierarchical in structure but is scored on testbench pass rate alone. Our contribution is the reuse-aware metric, the gold labels, and the multi-LLM eval protocol applied to existing benchmarks, rather than a new training pipeline.

VI. CONCLUSION

A reuse-aware metric separates LLMs that scratch-only benchmarks group together. The planner is the bottleneck of the system we built: holding the router and corpus fixed across 15 models, an F1 spread of 0.27 lives entirely in subblock decomposition and search-query specificity. The most important follow-up is the test-feedback retry loop on the integrator, where corpus-IP-stitching errors today land outside the planner’s awareness. We have also begun multi-seed (n=3) averaging on the four headline models to put proper variance bars on the leaderboard. Code, corpus, gold labels, and run scripts are open at <https://github.com/MattKotzbauer/rtl-mosaic>.

Self-assessment. We met the project goal of producing a working harness and a defensible benchmark. We did not show the harness end-to-end beating scratch generation on the `cpu_ip` set, where both are stuck at 1/9, so the routing-quality finding is the one we stand behind. Given another semester, we would build the test-feedback retry loop and run multi-seed averaging to put variance bars on the leaderboard.

TEAM CONTRIBUTIONS

Matt Kotzbauer: 30-IP corpus and self-tests, multi-provider runner, gold labels, eval metrics, figure pipeline, and the report. **Warren Zhu**: planner, integrator, ChipBench baseline runner, MCP IP-router, end-to-end harness CLI.

REFERENCES

- [1] Semico Research, “Semiconductor IP market: SoC share by area,” Industry report, 2024.
- [2] SiliconMind-V1: “Multi-agent distillation and debug for Verilog generation,” *arXiv preprint arXiv:2603.08719*, 2026.

- [3] ChipBench: “A realistic-module benchmark for LLM-driven RTL design,” *arXiv preprint* arXiv:2601.21448, 2026.
- [4] M. Liu et al., “VerilogEval: Evaluating large language models for Verilog code generation,” NVIDIA, *arXiv* arXiv:2309.07544, 2023.
- [5] K. Chang et al., “ChipGPT: How far are we from natural language hardware design,” *arXiv* arXiv:2305.14019, 2023.
- [6] S. Thakur et al., “VeriGen: A large language model for Verilog code generation,” *arXiv* arXiv:2308.00708, 2023.
- [7] M. Liu et al., “ChipNeMo: Domain-adapted LLMs for chip design,” NVIDIA, *arXiv* arXiv:2311.00176, 2023.
- [8] “MAGE: A multi-agent engine for automated RTL design,” *arXiv* arXiv:2412.07822, 2024.